

Part 10. Tree-based models - part 2

Dr. Nguyen Quang Huy

July 22, 2020

Tree-based model

- Tree-based model is a low-bias and high-variance method
- In this section, we introduce variance-reduction techniques in tree-based model including *bagging*, *random forests*, and *boosting*.
- These approaches reduce the variance of tree-based models by producing multiple trees which are then combined to yield a single prediction.
- The improvements in prediction accuracy of these approaches are at the expense of some loss in model interpretation.
- The topic to be discuss in this section
 - Bootstrapping
 - Bagging and random forest
 - Boosting

Bootstrapping

Bootstrap is a widely applicable and extremely powerful statistical tool that can be used to quantify the uncertainty associated with a given estimator.

```
summary(lm(medv~lstat,data=Boston))
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat, data = Boston)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

##	-15.168	-3.990	-1.318	2.034	24.500
----	---------	--------	--------	-------	--------

```
##
```


```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

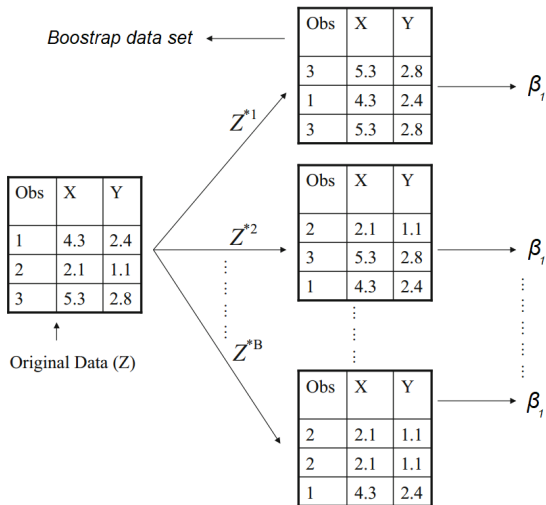
## (Intercept)	34.55384	0.56263	61.41	<2e-16 ***
## lstat	-0.95005	0.03873	-24.53	<2e-16 ***

Boostrapping

```
N<-1000
beta2<-rep(0,N)
for (i in 1:N){
  x<-sample(1:nrow(Boston),nrow(Boston),replace=TRUE)
  beta2[i]<-lm(medv~lstat,data=Boston[x,])$coef[2]
}
plot(density(beta2))
```



Boostrapping



Bootstrap is useful in situations in which it is hard to compute the standard deviation of an estimator

Bagging

- Given a set of n independent observations Z_1, \dots, Z_n , each with variance σ^2 , the variance of the mean \bar{Z} of the observations is given by σ^2/n .
- If we have n independent training data sets, we could calculate $\hat{f}_1, \dots, \hat{f}_n$ and average them in order to obtain a single low-variance model.
- It is not practical because we do not have access to multiple training sets. Instead, we can bootstrap, by taking repeated samples from the (single) training data set.
- Suppose that B different bootstrapped training data sets, we can train our method on the b^{th} bootstrapped training set in order to get $f^b(x)$, and finally average all the predictions

$$\hat{f}(x) = \frac{1}{B} \sum_{b=1}^B f^b(x)$$

Bagging

To apply bagging to regression trees,

- Construct B regression trees using B bootstrapped training sets.
- These trees are grown deep and are not pruned.
- Each individual tree has high variance, but low bias.
- Averaging these B trees to obtain prediction with lower variance than single tree

To apply bagging to classification trees, we can record the class predicted by each of the B trees, and take a majority vote: the overall prediction is the most commonly occurring majority class among the B predictions.

Bagging and Boston data set

```
library(randomForest)
dat<-Boston
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}
for (col in names(dat)){
  if((col!="medv")&class(dat[,col]) %in% c("integer","numeric")){
    dat[,col]<-standardize(dat[,col])
  }
}
set.seed(1)
test_index<-createDataPartition(dat$medv, times = 1, p = 0.2, sampleWithReplacement=FALSE)
train<-dat[-test_index,]
test<-dat[test_index,]
```


Bagging and Boston data set


```
bag.fit<-randomForest(medv~.,  
                      data=train,  
                      mtry=ncol(train)-1, #number of predictors  
                      importance=TRUE,  
                      ntree=200) # number of tree  
bag.pred<-predict(bag.fit,newdata=test)  
sqrt(mean((bag.pred-test$medv)^2))
```

```
## [1] 2.942067
```

Increase the number of tree and observe the test error ?

Bagging and Boston data set

Bagging – number of tree and test error



Bagging

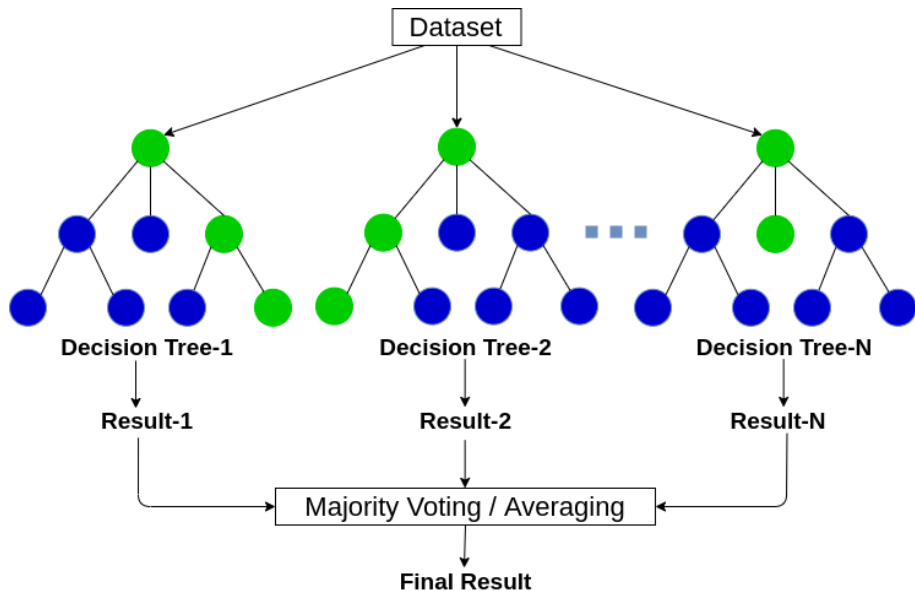
- Bagging results in improved accuracy over prediction using a single tree.
- Unfortunately, it can be difficult to interpret the resulting model.
- Bagging improves prediction accuracy at the expense of interpretability.
- We can obtain an overall summary of the importance of each predictor using the RSS (for bagging regression trees) or the Entropy index (for bagging classification trees)

##	%IncMSE	IncNodePurity
## crim	21.346694	999.88054
## zn	2.203785	30.26316
## indus	12.217769	231.67418
## chas	0.418869	39.55540
## nox	32.784004	838.88029
## rm	68.741107	14131.42916

Random forest

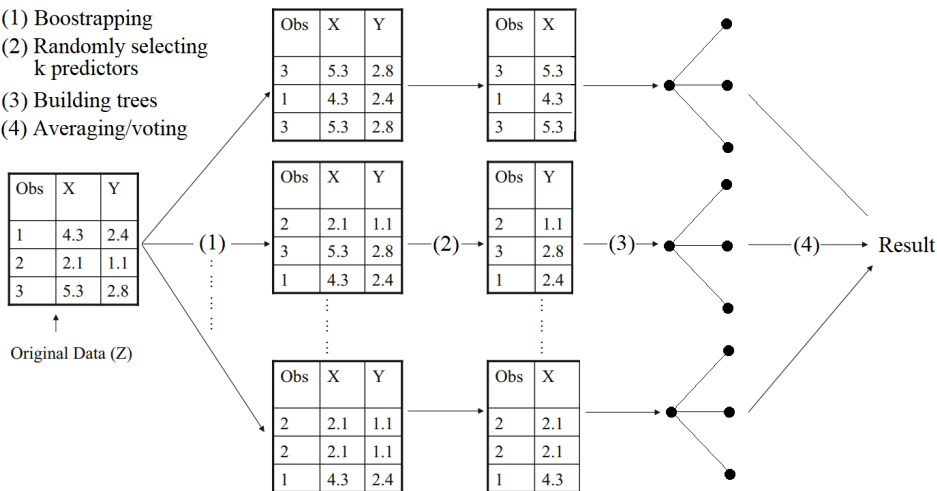
- Random forests provide an improvement over bagged trees by way of a random small tweak that *decorrelates* the trees
- If there is one very strong predictor in the data set, along with a number of other moderately strong predictors, all of the trees in bagging will use the strongest predictor in the top split.
- The predictions from the bagged trees will be highly correlated.
- Averaging many highly correlated responses does not lead to as large of a reduction in variance.
- Random forests overcome this problem by forcing each split to consider only a subset of the predictors (m random predictors instead of all p predictors)
- How to choose value of m , the number of random predictors using in each split, is out of scope of this course. In practice, they use $m = \lfloor \sqrt{p} \rfloor$ in classification and $m = \lfloor p/3 \rfloor$ in regression

Random forest




Random forest

- (1) Bootstrapping
- (2) Randomly selecting k predictors
- (3) Building trees
- (4) Averaging/voting



Random forest and Boston data set



Random forest and Default data set

Using bagging and random forest to predict *default* variables.

```
dat<-Default
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}
for (col in names(dat)){
  if((col!="default")&class(dat[,col]) %in% c("integer","numeric")){
    dat[,col]<-standardize(dat[,col])
  }
}
set.seed(1)
test_index<-createDataPartition(dat$default, times = 1, p = 0.2)
train<-dat[-test_index,]
test<-dat[test_index,]
```


Random forest and Default data set

Using bagging and random forest to predict *default* variables.

```
bag.fit<-randomForest(default~.,  
                        data=train,mtry=3, #number of predictors  
                        importance=TRUE,ntree=500) # number off  
bag.pred<-predict(bag.fit,newdata=test,type="class")  
table(bag.pred,test$default)
```

```
##  
## bag.pred    No   Yes  
##           No 4788 114  
##           Yes  46  53
```

Boosting

- Boosting is a general approach that can be applied to many statistical learning methods for regression or classification.
- In random forest method, the trees are grown independently using multiple copies of the original training data set using the bootstrap.
- Applying boosting to tree-based model, the trees are grown *sequentially*: each tree is grown using information from previously grown trees
- Instead of growing some large trees which fit the data hard and potentially overfitting, the boosting approach *learns slowly*
- Boosting involves combining a large number of decision trees, f^1, \dots, f^B where f^{j+1} is the tree grown from fitting predictors to the residuals from f^j .

Boosting for regression tree


Boosting with parameter B , d , λ .

- Step 1: Let $\hat{f} = 0$ and $res_i = y_i$ for all i in the training set.
- Step 2: For each $b = 1, 2, \dots, B$
 - Fit a d -nodes tree to (\mathbf{X}, res) and get a prediction \hat{f}_b
 - Update the prediction: $\hat{f} = \hat{f} + \lambda \times \hat{f}_b$
 - Update the residual: $res = res - \lambda \times \hat{f}_b$
- Step 3: Final output of boosted model:

$$Output = \sum_{b=1}^B \lambda \hat{f}_b$$

where B is the number of trees and λ is the learning rate. (Typical values of λ are 0.01 and 0.001)

Boosting



Boosting and Boston data set

```
library(gbm)
boost.fit<-gbm(medv~.,data=train,
               n.trees = 5000, # parameter B
               interaction.depth=5, # parameter d
               shrinkage = 0.005) # parameter lambda

## Distribution not specified, assuming gaussian ...

boost.pred<-predict(boost.fit, newdata=test, n.trees=5000)
sqrt(mean((boost.pred-test$medv)^2))

## [1] 2.755437
```

Xgboost

```
library(xgboost)
dtrain<-data.matrix(dplyr::select(train,-medv))
xgb<-xgboost(data = dtrain,
              label = train$medv,nround=1000,
              max.depth = 3, eta = 0.05,verbose = 0)

pred <- predict(xgb, data.matrix(dplyr::select(test,-medv)))
sqrt(mean((pred-test$medv)^2))

## [1] 2.572969
```