

Part 9 Linear model extensions

Dr. Nguyen Quang Huy

July 20, 2020

Overview

In this section, we will discuss about the following topic:

- ➊ Regression splines
 - Piecewise polynomials
 - Constraints and splines
 - Discussion on the number and the locations of the knots
- ➋ Smoothing splines
 - Introduction to smoothing splines
 - Choosing the smoothing parameter
- ➌ Local regression
- ➍ Generalized additive models (GAM)
- ➎ Generalized linear models (GLM)

Regression splines

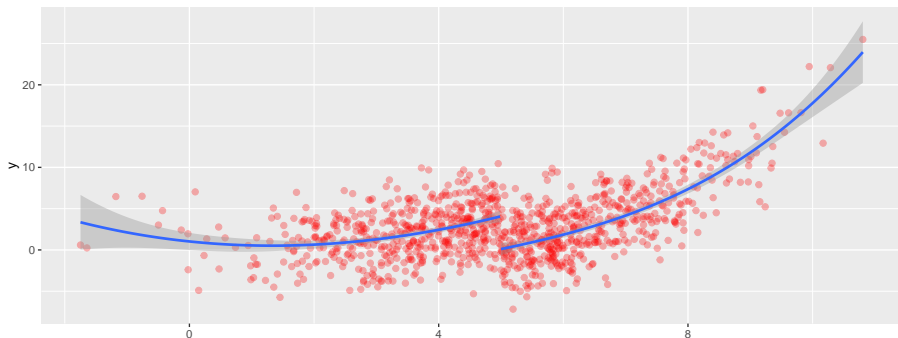
- Instead of fitting a high-degree polynomial over the entire range of \mathbf{X} , piecewise polynomial regression fits separate low-degree polynomials over different regions of \mathbf{X} (splines)
- A piecewise cubic polynomial with 1 knot at points k_1 takes the form

$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x < k_1 \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & \text{if } k_1 \leq x \end{cases}$$

- In general, if we place J different knots throughout the range of X , then we will fit $(J + 1)$ different polynomials.
- **Problem:** the function is discontinuous at knots.

Regression splines

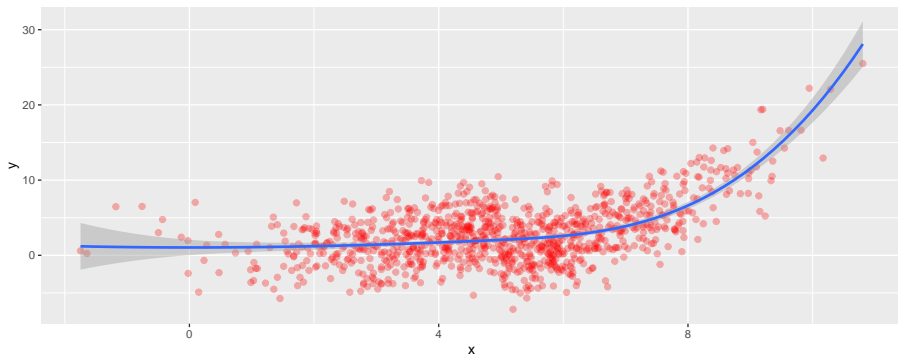
```
x<-rnorm(1000,5,2)
y=ifelse(x<5,0.03*x^3-0.2*x+1,0.02*x^3-0.1*x-2)+rnorm(1000,0,3)
z<-ifelse(x<5,TRUE,FALSE)
dat<-data.frame(x,y,z)
dat%>%ggplot(aes(x,y,group=z))+geom_point(col="red",cex=2,alpha=0.5)
  geom_smooth(method=lm,formula=y~poly(x,3,raw=TRUE))
```



Regression splines

To solve the uncontinuous problem, they add two additional constraints: both the first and second derivatives of the piecewise polynomials are continuous.

```
dat%>%ggplot(aes(x,y))+geom_point(aes(group=z),col="red",cex=2)  
  geom_smooth(method=lm,formula=y~bs(x,knots=5))
```



Regression splines

We can prove that, the estimation of a **continuous cubic splines** with m knots, k_1, k_2, \dots, k_m is similar to the estimation of

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \beta_4 h_1(x_i) + \beta_5 h_2(x_i) + \dots + \beta_{m+3} h_m(x_i) + \epsilon_i$$

where

$$h_j(x_i) = \begin{cases} (x - k_j)^3 & \text{if } x \geq k_j \\ 0 & \text{if } x < k_j \end{cases}$$

Where should we place the knots

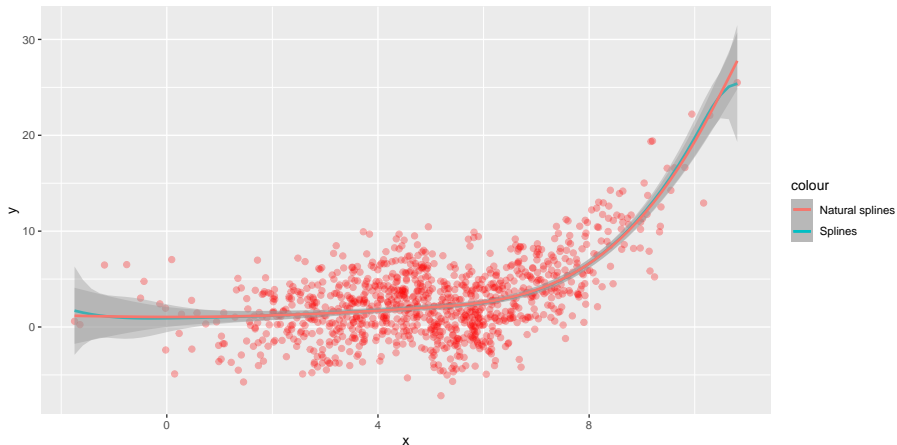
- Manually: place more knots in places where the function might vary most rapidly, and to place fewer knots where it seems more stable.
- Automatically: specify the desired degrees of freedom (number of β), and have the software automatically place the corresponding number of knots (cross validation)

Regression splines

- A (continuous) cubic splines has $(m + 3)$ parameters where m is the number of knots
- Splines can have high variance at the outer range of the predictors, that is, when X takes on either a very small or very large value
- They introduce “natural spline” which is a spline with additional boundary constraints: the function is required to be **linear** at the boundary
 - In the region where X is spline smaller than the smallest knot
 - or in the region where X larger than the largest knot
- A natural cubic splines has m parameteres where m is the number of knows

Regression splines

```
dat%>%ggplot(aes(x,y))+geom_point(aes(group=z),col="red",cex=2)  
  geom_smooth(method=lm,formula=y~bs(x,knots=c(0,5,10)),aes(col=group))  
  geom_smooth(method=lm,formula=y~ns(x,knots=c(0,5,10)),aes(col=group))
```



Smoothing splines

- What we do is find some function $f(x)$ that fits the observed data well i.e. minimize $RSS = \sum_{i=1}^n (y_i - f(x_i))^2$.
- If we don't put any constraints on $f(x)$, then we can make RSS small simply by choosing f such that it interpolates all of the y_i , or it is easy to overfit
- What we want function f that makes RSS small, but f is also **smooth**.
- We find the function f that minimizes

$$\sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \int f''(x) dx$$

where $\lambda \int f''(x) dx$ is a penalty term.

- The second derivative is the speed of change of the slope i.e. associated with the roughness of the curve.

Smoothing splines

- It can be shown that this problem has an explicit, finite-dimensional, which is a **natural cubic spline** with knots at the unique values of the $x_i, i = 1, \dots, n$
- It seems that function f is over-parametrized (n degree of freedom)
 - Regression splines with n knots has $(n + 4)$ parameter
 - Linear constraints when $x < \min(k_i)$ and $x > \max(k_i) \rightarrow n$ parameter left
- The penalty term translates to a penalty on the spline coefficients, which are shrunk the model toward the linear fit.
- The solution has the following form, where b_i are basis functions

$$f(x) = \sum_{i=1}^n \beta_i b_i(x)$$

Smoothing splines

Solution of smoothing splines (forget it :D): for each λ , find vector β to minimize

$$(\mathbf{y} - \mathbf{B}\beta)' (\mathbf{y} - \mathbf{B}\beta) + \lambda \beta' \Omega \beta$$

where $B_{ij} = b_i(x_j)$ and $\Omega_{ij} = \int b_i''(t)b_j''(t)$. The solution is

$$\hat{\beta} = (B' B + \lambda \Omega)^{-1} B' \mathbf{y}$$

Thus, the smoothing splines is

$$\hat{f}(x) = \sum_{i=1}^n \hat{\beta}_i b_i(x)$$

Let $\hat{\mathbf{f}} = (f(x_1), f(x_2), \dots, f(x_n))$, we have

$$\hat{\mathbf{f}} = \mathbf{B}\hat{\beta} = \mathbf{B} (B' B + \lambda \Omega)^{-1} B' \mathbf{y} = \mathbf{S}_{\lambda} \mathbf{y}$$

Smoothing splines

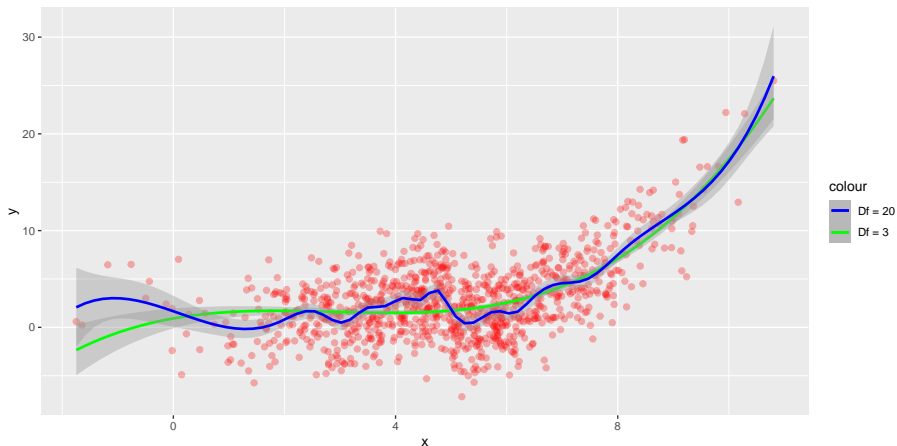
The effective degree of freedom is defined as “the sum of the sensitivities of the fitted values with respect to the observed response values”

$$\sum_{i=1}^n \frac{\partial \hat{y}_i}{\partial y_i} = \text{trace}(\mathbf{S}_\lambda)$$

- In a linear regression model with k predictors, the degree of freedom is the number of parameters: $(k+1)$
- In smoothing splines, there are n parameters, but they are constrained by others (because of λ) \rightarrow the effective degree of freedom is $\text{trace}(\mathbf{S}_\lambda)$
- Larger effective degree of freedom, higher variance.
- Lower effective degree of freedom, lower variance.

Regression splines

```
dat%>%ggplot(aes(x,y))+geom_point(aes(group=z),col="red",cex=2)  
  geom_smooth(method=lm,formula=y~ splines::bs(x,df=3) ,aes(colour=group))  
  geom_smooth(method=lm,formula=y~ splines::bs(x,df=20),aes(colour=group))
```



Local regression

Local regression is a different approach for fitting flexible non-linear functions, which involves computing the fit at a target point using **only the regression nearby training observations**.

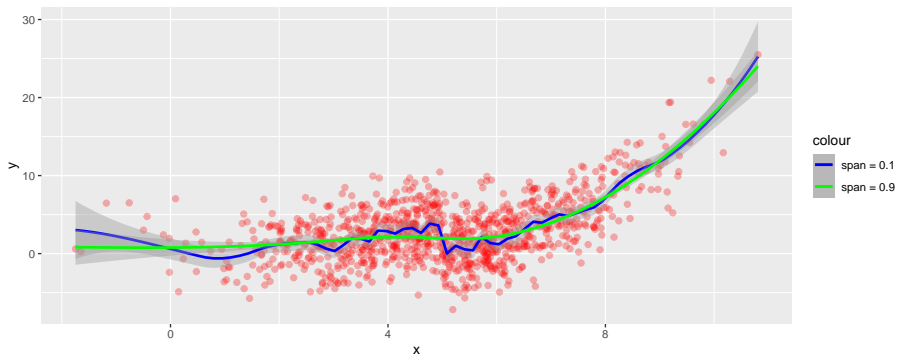
- Gather the fraction $s = k/n$ (span) of training points whose x_i are closest to x .
- Choose a weight function K to each point in this neighborhood so that the point furthest from x has weight zero, and the closest has the highest weight; $K_i = K(x_i, x)$
- Fit a weighted least squares regression: finding β_0, β_1 that minimize

$$\sum K_i(y_i - \beta_0 - \beta_1 x_i)^2$$

- Return $f(x) = \hat{\beta}_0 + \hat{\beta}_1 x$

Regression splines

```
dat%>%ggplot(aes(x,y))+geom_point(aes(group=z),col="red",cex=2)  
  geom_smooth(method="loess",span=0.1,aes(colour="span = 0.1"))  
  geom_smooth(method="loess",span=0.9,aes(colour="span = 0.9"))  
  scale_colour_manual(values=c("blue","green"))
```



Generalized additive models

- We have presented a number of approaches for flexibly predicting a response Y on the basis of a single predictor X
- **Generalized additive models (GAMs)** provide a general framework for extending a standard linear model by allowing non-linear functions of each of the variables.
- GAM can be write as

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i$$

where f_i can be a constant, a polinomal, a natural splines, or a smoothing splines ...

- When f_j is a smoothing splines, the least square method can not be used. GAM method fits a model involving multiple predictors by repeatedly updating the fit for each predictor in turn, that is, apply the fitting method for that variable to a *partial residual*

Generalized additive models

Pros of GAMs

- 1 GAMs allow us to automatically fit a non-linear f_j to each X_j , we do not need to manually try out many different transformations on each variable individually.
- 2 The non-linear function f_j can potentially make more accurate predictions.
- 3 We can examine the effect of each X_j on Y individually. If we are interested in inference, GAMs provide a useful representation.

Cons of GAMs

- 1 The main limitation of GAMs is that the model is restricted to be additive (linears). With many variables, important interactions can be missed.

Generalized additive models - practice

Load dataset **Boston** from **MASS** packages and build GAM to predict **medv**.

```
dat<-Boston
```

- 1 Standardize all numerical variables (except for **medv**) and split data into training set and test set (80%-20%)
- 2 Build a linear model when $medv \sim lstat$ (use polynomial, splines, natural splines, smoothing splines). Which model has the lowest error (on test dataset)?
- 3 Build a GAM where $medv$ depends on all predictors.

Generalized additive models - practice

```
dat<-Boston
# STANDARDIZE
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}
for (col in names(dat)){
  if((col!="medv")&class(dat[,col]) %in% c("integer","numeric")){
    dat[,col]<-standardize(dat[,col])
  }
}
# SPLITTING INTO TRAIN - TEST
set.seed(1)
test_index<-createDataPartition(dat$medv, times = 1, p = 0.2, sampleWithReplacement=FALSE)
train<-dat[-test_index,]
test<-dat[test_index,]
```

Generalized additive models - practice

```
## POLINOMIAL
```

```
poly.fit<-lm(medv~poly(lstat,4,raw=TRUE),data=train)  
medv.pred<-predict(poly.fit,test)  
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 5.61816
```

```
## CUBIC SPLINES (CONTINUOUS)
```

```
cubic.splines<-lm(medv~bs(lstat,knots = c(0)),data=train)  
medv.pred<-predict(cubic.splines,test)  
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 5.62075
```

Generalized additive models - practice

```
library(gam)
# NATURAL SPLINE (LINEAR when x small and x large)
natural.splines<-lm(medv~ns(lstat,df=5),data=train)
medv.pred<-predict(natural.splines,test)
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 5.626698
```

```
# SMOOTHING SPLINES (require gam package)
df1<-smooth.spline(train$medv,train$lstat,cv=TRUE)$df
smoothing.spl<-gam(medv~s(lstat,df=df1),data=train)
medv.pred<-predict(smoothing.spl, newdata = test)
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 5.606942
```

Generalized additive models - practice

```
# Different functions in GAM model
gam1<-gam(medv~s(lstat,6)+lo(dis,0.3)+
          poly(crim,4,raw=TRUE)+zn,data=train)
medv.pred<-predict(gam1, newdata = test)
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 5.055312
```

```
# ADDITIVE
gam2<-gam(medv~s(lstat)+s(dis)+
          s(crim)+s(zn),data=train)
medv.pred<-predict(gam2, newdata = test)
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 4.957187
```

Generalized additive models - practice

```
## MANUALLY CHOOSE MODEL :)
```

```
gam3<-gam(medv~s(lstat)+s(crim)+s(zn)+s(indus)+  
          s(nox)+s(rm)+s(dis)+s(age)+s(tax)+  
          s(ptratio)+chas, data=train)  
medv.pred<-predict(gam3, newdata = test)  
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 3.769918
```

```
gam4<-mgcv::gam(medv~s(lstat)+s(crim)+s(zn)+  
                s(indus)+s(nox)+s(rm)+s(dis)+  
                s(age)+s(tax)+s(ptratio)+chas+s(black)+rad,  
                data=train,select=TRUE)  
medv.pred<-predict(gam4, newdata = test)  
sqrt(mean((medv.pred-test$medv)^2))
```

```
## [1] 3.299333
```

Generalized additive models

GAMs are generally fit using a *backfitting* approach.

```
n<-1000
set.seed(1)
x1<-rnorm(n,0,1)
x2<-rnorm(n,0,1)
x3<-rnorm(n,0,1)
y<-1+2*x1+3*x2+4*x3+rnorm(n,0,5)
lm(y~x1+x2+x3) # easy to perform a multi-regression
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ x1 + x2 + x3)
```

```
##
```

```
## Coefficients:
```

## (Intercept)	x1	x2	x3
## 1.085	2.115	3.076	4.046

Generalized additive models

Suppose that you only have a computer to perform simple linear regression.

- 1 Fit the model $Y \sim X_1$ to obtain $\hat{\beta}_1$.
- 2 Fit the model $Y - \hat{\beta}_1 X_1 \sim X_2$ to obtain $\hat{\beta}_2$.
- 3 Fit the model $Y - \hat{\beta}_1 X_1 - \hat{\beta}_2 X_2 \sim X_3$ to obtain $\hat{\beta}_3$.
- 4 Back to step (1), replace Y by $Y - \hat{\beta}_2 X_2 - \hat{\beta}_3 X_3$.
- 5 Write a for loop to repeat these steps N times ($N = 10$)
- 6 With $\hat{\beta}_1$, $\hat{\beta}_2$ and $\hat{\beta}_3$ from step (5), calculate $\hat{\beta}_0$ as follow

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}_1 - \hat{\beta}_2 \bar{X}_2 - \hat{\beta}_3 \bar{X}_3$$

Generalized additive models

```
N<-5
b<-matrix(0,3,N)
for(i in 2:N){
  b[1,i]<-lm(y-b[2,i-1]*x2-b[3,i-1]*x3~x1)$coef[2]
  b[2,i]<-lm(y-b[1,i-1]*x1-b[3,i-1]*x3~x2)$coef[2]
  b[3,i]<-lm(y-b[2,i-1]*x2-b[1,i-1]*x1~x3)$coef[2]
}
b[,N]

## [1] 2.114983 3.075799 4.046309

b[1,]

## [1] 0.000000 2.333841 2.105754 2.115697 2.114983

mean(y)-b[1,N]*mean(x1)-b[2,N]*mean(x2)-b[3,N]*mean(x3)

## [1] 1.085474
```

Generalized linear models (GLM)

We have the linear model

$$Y = \beta_1 X_1 + \cdots \beta_p X_p + \epsilon = \mathbf{X}'\beta + \epsilon$$

with $\epsilon \sim \mathcal{N}(0, \sigma^2)$. Or we can write

$$\mathbb{E}(Y|X) = \mathbf{X}'\beta$$

where $Y|X$ has normal distribution function. Consider a more general case, where

$$g(\mathbb{E}(Y|X)) = \mathbf{X}'\beta$$

where g is called link function and $Y|X$ has distribution function F . g must be monotone and differentiable.

Generalized linear models

- In linear model, we assume that $Y|X = x_i$ has normal distribution, but it is not appropriate in many situation.
- In credit risk modeling, Y is binary, $\rightarrow Y|X = x_i$ is binary, we can not assume that $Y|X = x_i$ has normal distribution. A suitable distribution for $Y|X = x_i$ is Bernoulli distribution.
- When Y is a counting number, a number of claim in a year for example, $Y|X = x_i$ is a counting number. $Y|X = x_i$ can not have normal distribution. $Y|X = x_i$ can be a Poisson or a Negative Binomial random variable.
- Even when Y is continuous, but $Y > 0$, then $Y|X = x_i > 0 \forall x_i$; the normality assumption of $Y|X = x_i$ should be considered.

Generalized linear models

- In linear regression, we try to build model where $\mathbb{E}(Y|X = x_i)$ is a linear function of x_i i.e. find β_0 and β_1 such that $\mathbb{E}(Y|X = x_i) = \beta_0 + \beta_1 x_i$.
- $\beta_0 + \beta_1 x_i$ can take any value on \mathbb{R}
- $\mathbb{E}(Y|X = x_i)$, however, depends on the distribution of $Y|X$.
- If $Y|X = x_i$ has binomial distribution then $\mathbb{E}(Y|X = x_i) \in [0, 1]$.
- If $Y|X = x_i$ has poisson distribution then $\mathbb{E}(Y|X = x_i) \in [0, \infty)$.
- To match the range of $\mathbb{E}(Y|X = x_i)$ to set \mathbb{R} , they introduce a link function g .
 - For example, $\mathbb{E}(Y|X = x_i) \in [0, 1]$, we need function g such that $g : [0, 1] \rightarrow \mathbb{R}$
 - When $\mathbb{E}(Y|X = x_i) \in [0, \infty)$, we need function g such that $g : [0, \infty] \rightarrow \mathbb{R}$

Generalized linear models

- When $Y|X = x_i$ is a Bernoulli random variable, any function $g : [0, 1] \rightarrow \mathbb{R}$ can be a link function.
- However, when $Y|X = x_i \sim \mathcal{B}(p_i)$, the logit function is the canonical link function. The logit function is defined as follows

$$g(x) = \text{logit}(x) = \log \left(\frac{x}{1-x} \right)$$

- Definition of canonical link function is out-of-scope of this course. All you should know is that when the link function g is the canonical link, we have an analytic solution for β_0, β_1, \dots
- It explains why when $Y|X = x_i$ is a Bernoulli, they choose g as logit function. For this choice of Y and g , we often call **logistic regression**,

Generalized linear models

- Quick question 1: Name two other functions that can be a link function when $Y|X = x_i$ is a Bernoulli

Generalized linear models

- Quick question 1: Name two other functions that can be a link function when $Y|X = x_i$ is a Bernoulli
- $\Phi^{-1}(x)$: the inverse function of the distribution function of standard normal random variable (Probit regression)

Generalized linear models

- Quick question 1: Name two other functions that can be a link function when $Y|X = x_i$ is a Bernoulli
- $\Phi^{-1}(x)$: the inverse function of the distribution function of standard normal random variable (Probit regression)
- $\log(-\log(1 - x))$

Generalized linear models

- Quick question 1: Name two other functions that can be a link function when $Y|X = x_i$ is a Bernoulli
- $\Phi^{-1}(x)$: the inverse function of the distribution function of standard normal random variable (Probit regression)
- $\log(-\log(1 - x))$
- Quick question 2: Name a function that can be a link function when $Y|X = x_i$ is a Poisson

Generalized linear models

- Quick question 1: Name two other functions that can be a link function when $Y|X = x_i$ is a Bernoulli
- $\Phi^{-1}(x)$: the inverse function of the distribution function of standard normal random variable (Probit regression)
- $\log(-\log(1 - x))$
- Quick question 2: Name a function that can be a link function when $Y|X = x_i$ is a Poisson
- $\log(x)$

Generalized linear models - estimation

- Suppose that $Y|(X = x_i)$ has density function/p.m.f f_i with parameter θ_i
- f_i is the density/p.m.f function of random variable $Y_i = Y|(X = x_i)$ where $g(\mathbb{E}(Y_i)) = \mathbf{x}_i'\beta$
- Because $\mathbb{E}(Y_i)$ is a function of θ_i , we have can write $\theta_i = h(\mathbf{x}_i'\beta)$
- We must find β to maximize the log likelihood function

$$L(\beta) = \text{Log} \prod_{i=1}^n f_i(y_i) = \sum_{i=1}^n \text{Log}(f_i(y_i))$$

where paramater of f_i is $\theta_i = h(\mathbf{x}_i'\beta)$

Generalized linear models - estimation

- For example, when f_i is the p.m.f of Bernoulli random variable with parameter p_i and g is logit function: $g(x) = \log(x/(1-x))$
- What is the p.m.f f_i at y_i : $f_i(y_i) = \mathbb{P}(Y_i = y_i)$

Generalized linear models - estimation

- For example, when f_i is the p.m.f of Bernoulli random variable with parameter p_i and g is logit function: $g(x) = \log(x/(1-x))$

- What is the p.m.f f_i at y_i : $f_i(y_i) = \mathbb{P}(Y_i = y_i)$

$$f_i(y_i) = p_i^{y_i} \times (1 - p_i)^{(1-y_i)} \rightarrow \log(f_i) = y_i \log(p_i) + (1 - y_i) \log(1 - p_i)$$

- What is the relation between p_i and $\mathbf{x}_i' \beta$

$$g(\mathbb{E}(Y_i)) = \mathbf{x}_i' \beta$$

$$\rightarrow \log\left(\frac{p_i}{1 - p_i}\right) = \mathbf{x}_i' \beta$$

$$\rightarrow p_i = \frac{\exp(\mathbf{x}_i' \beta)}{1 + \exp(\mathbf{x}_i' \beta)}$$

Generalized linear models - estimation

- Write $\sum \log(f_i)$ as a function of β

$$\begin{aligned} L(\beta) &= \sum_{i=1}^n \log(f_i) \\ &= \sum_{i=1}^n y_i \log \left[\frac{\exp(\mathbf{x}_i' \beta)}{1 + \exp(\mathbf{x}_i' \beta)} \right] + (1 - y_i) \log \left[\frac{1}{1 + \exp(\mathbf{x}_i' \beta)} \right] \\ &= \sum_{i=1}^n y_i \mathbf{x}_i' \beta - \log \left[1 + \exp(\mathbf{x}_i' \beta) \right] \end{aligned}$$

- Methods to solve for β
 - Newton–Raphson method
 - Iteratively reweighted least squares method

Generalized linear models

- Using GLM model to predict credit card default (data **Default** in **ISLR** package)

```
dat<-Default  
head(Default)
```

```
##      default student   balance   income  
## 1         No      No  729.5265 44361.625  
## 2         No     Yes  817.1804 12106.135  
## 3         No      No 1073.5492 31767.139  
## 4         No      No  529.2506 35704.494  
## 5         No      No  785.6559 38463.496  
## 6         No     Yes  919.5885  7491.559
```

```
summary(Default$default)
```

```
##      No   Yes  
## 9667  333
```


Generalized linear models

```
# STANDARDIZE
```

```
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}  
for (col in names(dat)){  
  if(class(dat[,col]) %in% c("integer","numeric")){  
    dat[,col]<-standardize(dat[,col])  
  }  
}
```

```
# SPLITTING INTO TRAIN - TEST
```

```
set.seed(1)  
test_index<-createDataPartition(dat$default, times = 1, p = 0.3)  
train<-dat[-test_index,]  
test<-dat[test_index,]  
summary(train$default)
```

```
##      No   Yes
```

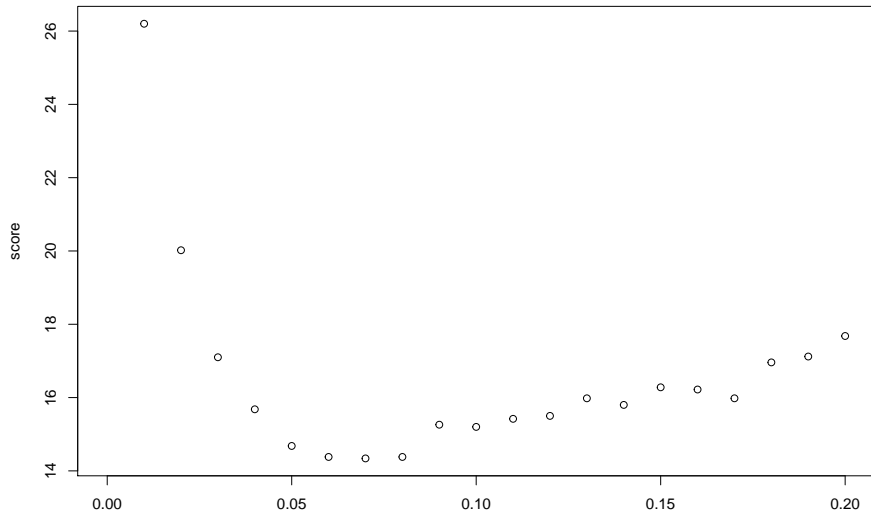
```
## 4833  166
```

Generalized linear models

```
logit<-glm(default~.,  
            family = binomial(link = "logit"),data=train)  
pred<-predict(logit,test,type="response")  
#Fixed cut point  
log.pred1<-as.factor(ifelse(pred>0.1,"Yes","No"))  
table(log.pred1,test$default)
```

```
##  
## log.pred1    No   Yes  
##           No  4563   45  
##           Yes  271  122
```

Generalized linear models



Generalized linear models

```
finalcut<-cutpoint[which.min(score)]
logit<-glm(default~.,
            family = binomial(link = "logit"),data=train)
pred<-predict(logit,test,type="response")
log.pred1<-as.factor(ifelse(pred>finalcut,"Yes","No"))
table(log.pred1,test$default)
```

```
##
## log.pred1    No    Yes
##           No  4465   35
##           Yes  369  132
```

Generalized linear models

```
finalcut<-cutpoint[which.min(score)]
probit<-glm(default~.,
             family = binomial(link = "probit"),data=train)
pred<-predict(probit,test,type="response")
log.pred1<-as.factor(ifelse(pred>finalcut,"Yes","No"))
table(log.pred1,test$default)
```

```
##
## log.pred1    No   Yes
##           No  4414   31
##           Yes  420  136
```

GLM and GAM

- In the general GLM model, we have

$$g(\mathbb{E}(Y|X = x_i)) = \mathbf{x}_i' \beta$$

- We could extend the right hand side of this formula by the sum of function of x_{ij}

$$g(\mathbb{E}(Y|X = x_i)) = \sum_{j=1}^p f_j(x_{ij})$$

where f_j can be a polynomial, natural splines, smoothing splines ...

GLM and GAM

```
logit<-gam(default~ student + s(income) + s(balance) ,  
            family = binomial,data=train)  
pred<-predict(logit,test,type="response")  
log.pred1<-as.factor(ifelse(pred>finalcut,"Yes","No"))  
table(log.pred1,test$default)
```

```
##  
## log.pred1    No    Yes  
##           No  4434    34  
##           Yes   400   133
```

Exercise

Exercise 1 In this exercise, we will analyze the **Wage** data set.

- Perform polynomial regression to predict wage using age. Use cross-validation to select the optimal degree d for the polynomial.
- Using natural splines and smoothing splines to predict wage using age.
- Make a plot of the resulting polynomial fit to the data.

Exercise 2 This exercise relates to the **College** data set.

- Split the data into a training set and a test set. Using out-of-state tuition as the response and the other variables as the predictors, perform forward stepwise selection on the training set in order to identify a satisfactory model that uses just a subset of the predictors.
- Fit a GAM on the training data, using out-of-state tuition as the response and the features selected in the previous step as the predictors. Plot the results, and explain your findings.