


# Part 10. Tree-based models

Dr. Nguyen Quang Huy

May 16, 2020

# Tree-Based model



- Tree-based method is used for both regression and classification.
- Tree-based methods are simple and useful for interpretation.
- This method can not be compative with other algorithms in accuracy.
- Combine the tree-based method with the technique of bagging, random forests or boosting, they often result in dramatic improvements in prediction accuracy.

# Regression tree - Boston dataset

Applying tree model to predict value of **medv** in Boston dataset. We use variable **lstat** to build the tree:

- With a cut-point  $c_1$ , we divide training dataset into two part:

$$R_1 : \text{lstat} < c_1$$

$$R_2 : \text{lstat} \geq c_1$$

let  $\mu_1$  and  $\mu_2$  are the means of **medv** in  $R_1$  and  $R_2$ .

- We find  $c_1$  to minimize the following

$$\sum_{i \in R_1} (\text{medv}_i - \mu_1)^2 + \sum_{j \in R_2} (\text{medv}_j - \mu_2)^2$$

- We predict values of **medv** in test set as follows:
  - If **lstat**  $< c_1$  then **medv**  $= \mu_1$
  - If **lstat**  $\geq c_1$  then **medv**  $= \mu_2$

# Regression tree - Boston data set

- Load Boston data set, standardize the numerical variables, split data into training and test set (80%-20%);

```
dat<-Boston
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}
for (col in names(dat)){
  if((col!="medv")&class(dat[,col]) %in% c("integer","numeric")){
    dat[,col]<-standardize(dat[,col])
  }
}
set.seed(1)
test_index<-createDataPartition(dat$medv, times = 1, p = 0.2, sampleWithReplacement=FALSE)
train<-dat[-test_index,]
test<-dat[test_index,]
```

# Regression tree - Boston data set

- We will build a tree for  $medv$  based on  $lstat$  variable using training set. The first step is to find  $c_1$  such that

$$\mathbf{R}_1 = \mathbf{lstat} \leq c_1; \mathbf{R}_2 = \mathbf{lstat} > c_1$$

$$\mu_1 = \mathbb{E}(medv | \mathbf{R}_1); \mu_2 = \mathbb{E}(medv | \mathbf{R}_2)$$

$$\sum_{i \in R_1} (medv_i - \mu_1)^2 + \sum_{j \in R_2} (medv_j - \mu_2)^2 \rightarrow \min$$

# Regression tree - Boston data set

- We will build a tree for *medv* based on *lstat* variable using training set. The first step is to find  $c_1$  such that

$$\mathbf{R}_1 = \text{lstat} \leq c_1; \mathbf{R}_2 = \text{lstat} > c_1$$

$$\mu_1 = \mathbb{E}(\text{medv} | \mathbf{R}_1); \mu_2 = \mathbb{E}(\text{medv} | \mathbf{R}_2)$$


$$\sum_{i \in R_1} (\text{medv}_i - \mu_1)^2 + \sum_{j \in R_2} (\text{medv}_j - \mu_2)^2 \rightarrow \min$$

```
K<-1000
c1<-seq(min(train$lstat),max(train$lstat),length=K)
Error<-rep(0,length(c1))
for(i in 1:K){
  m1<-mean(train$medv[train$lstat<c1[i]])
  m2<-mean(train$medv[train$lstat>=c1[i]])
  Error[i]<-sqrt(sum((train$medv[train$lstat<c1[i]]-m1)^2)+
                  sum((train$medv[train$lstat>=c1[i]]-m2)^2))
}
```

# Regression tree - Boston data set

Find  $c_1$  that minimizes the error (on training dataset)

```
plot(c1, Error, type="l")
```



# Regression tree - Boston data set

Find the next node: using the following function

```
Tree.cut<-function(y,x,K){
  c1<-seq(min(x),max(x),length=K)
  Error<-rep(0,length(c1))
  for(i in 1:K){
    m1<-mean(y[x<=c1[i]])
    m2<-mean(y[x>=c1[i]])
    Error[i]<-sqrt(sum((y[x<=c1[i]]-m1)^2)+
                  sum((y[x>=c1[i]]-m2)^2))
  }
  result<-list()
  result$cutpoint<-c1[which.min(Error)]
  result$error<-min(Error)
  Tree.cut<-result}
```



# Regression tree - Boston data set

```
c1<-Tree.cut(train$medv,train$lstat,500)$cutpoint
train.R1<-filter(train,lstat<=c1)
train.R2<-filter(train,lstat>c1)
tree1<-Tree.cut(train.R1$medv,train.R1$lstat,500)
tree2<-Tree.cut(train.R2$medv,train.R2$lstat,500)
tree1$error
```

```
## [1] 74.03281
```

```
tree2$error
```

```
## [1] 75.6511
```

# Regression tree - Boston data set

```
c1<-Tree.cut(train$medv,train$lstat,500)$cutpoint
train.R1<-filter(train,lstat<=c1)
train.R2<-filter(train,lstat>c1)
tree1<-Tree.cut(train.R1$medv,train.R1$lstat,500)
tree2<-Tree.cut(train.R2$medv,train.R2$lstat,500)
tree1$error
```

```
## [1] 74.03281
```

```
tree2$error
```


```
## [1] 75.6511
```

```
ifelse(tree1$error<=tree2$error,
        tree1$cutpoint,tree2$cutpoint)
```

```
## [1] -1.147438
```

# Regression tree - Boston data set

```
## [1] -0.705836
```




This is a tree with 2 leaves and 1 node.

# Regression tree - Boston data set

Applying the same procedure of finding cut-point in each region  $R_1$  and  $R_2$ , we have a large tree

```
## [1] -0.705836
```

```
## [1] -1.147438
```




This is a tree with 3 leaves (3 predictions to medv) and 2 nodes

# Regression tree - Boston dataset

Function `tree()` in library **tree** fits the data set by tree-based model.

```
tree.fit <- tree(medv ~ lstat, data = train)
plot(tree.fit)
text(tree.fit, pretty = 0)
```



# Regression tree - Boston dataset

- The tree grows until a stopping criteria is reached.
- We could change these stopping criterias using *control* = *tree.control(...)* inside the tree function.
- There are 4 criterias inside the *tree.control*: *nobs*, *mincut* (5), *minsize* (10), *mindev* (0.01).
  - *nobs*: the number of observations you have available
  - *mincut*: the minimum number of observations each child node has to contain
  - *minize*: specifies the minimum number of observations a node must contain to be split
  - *mindev*: specifies the minimum ratio of the deviance at the root-node to the deviance at the node under consideration.

# Regression tree - Boston dataset

*# A LARGE TREE*

```
setup1<-tree.control(nobs=nrow(train),  
                     mincut=2,minsize=4,mindev=0.001)  
tree.fit.full<-tree(medv~lstat,data=train,  
                    control = setup1)  
plot(tree.fit.full)  
text(tree.fit.full,pretty=0)
```

*# A SMALL TREE*

```
setup2<-tree.control(nobs=nrow(train),  
                     mincut=250,minsize=500,mindev=0.001)  
tree.fit.full<-tree(medv~lstat,data=train,  
                    control = setup2)  
plot(tree.fit.full)  
text(tree.fit.full,pretty=0)
```


# Regression tree - tree pruning

- A large tree produce good predictions on the training set, but is likely to overfit the data (low bias but large variance).
- The such tree may produce poor performance on the test set in general.
- A smaller tree may result in lower variance at the cost of a little bias.
- Number of nodes (or leaves) in the smaller tree should be selected using cross validation approach.
- Function *prune.tree(treename, best = L)* is used to prune a large tree until number of leaves is  $L$



# Regression tree - tree pruning

```
tree.fit<-tree(medv~lstat,data=train)
prune.tree<-prune.tree(tree.fit,best = 5) # 5 leaves
plot(prune.tree)
text(prune.tree,pretty = 0)
```



# Regression tree - cross validation


- Function **cv.tree** can be used to perform cross-validation on training data set.

```
setup<-tree.control(nobs=nrow(train),  
                    mincut=2,minsize=4,mindev=0.005)  
tree.fit<-tree(medv~lstat,data=train,control = setup)  
cv.tree1<-cv.tree(tree.fit,K=5)  
L<-cv.tree1$size[which.min(cv.tree1$dev)]  
mytree<-prune.tree(tree.fit,best=L)  
tree.pred<-predict(mytree,newdata=test)  
sqrt(mean((tree.pred-test$medv)^2))  
  
## [1] 5.726087
```


# Regression tree - multiple variables

Similar to linear models, we can use multi-variables in a tree model.


```
tree.fit2<-tree(medv~lstat+rm,data = train)
prune.tree2<-prune.tree(tree.fit2,best = 3)
plot(prune.tree2)
text(prune.tree2,pretty=0)
```




# Regression tree - Boston dataset




# Regression tree - Boston dataset




# Regression tree - Boston dataset



# Regression tree - Boston dataset



# Regression tree - Boston dataset





# Regression tree

The process of building a regression tree

- Step 1: divide the set of possible values for  $X_1, X_2, \dots, X_p$  into  $J$  distinct and non-overlapping rectangles (hyper-rectangles in high dimension)  $R_1, R_2, \dots, R_J$  that minimize

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \bar{y}_{R_j})^2$$

- For every observation that falls into the region  $R_j$ , we make the same prediction, which is the mean of the response values for the training observations in  $R_j$  i.e.  $\bar{y}_{R_j}$

However, it is impossible to consider every possible partition of the feature space into  $J$  boxes. In practice, they take the recursive binary splitting algorithm

# Regression tree

## Recursive binary splitting algorithm

- is a top-down algorithm where it begins at the top of the tree and then successively splits the predictor space; each split is indicated via two new branches further down on the tree.
- is a greedy algorithm because at each step of the tree-building process, the best split is made at that particular step, rather than looking ahead and picking a split that will lead to a better tree in some future step.
- At each step, we select a predictor  $X_j$  and cutpoint  $c$  that minimize

$$\sum_{i \in R_1} (y_i - \bar{y}_{R_1})^2 + \sum_{k \in R_2} (y_k - \bar{y}_{R_2})^2$$

where  $R_1 = X | X_j \leq c$  and  $R_2 = X | X_j > c$

# Regression tree


```
setup<-tree.control(nobs=nrow(train),  
                    mincut=2,minsize=4,mindev=0.005)  
tree.fit<-tree(medv~lstat+rm,data=train,control = setup)  
cv.tree1<-cv.tree(tree.fit,K=5)  
L<-cv.tree1$size[which.min(cv.tree1$dev)]  
# number of leaves by cross validation  
L
```

```
## [1] 11
```


```
mytree<-prune.tree(tree.fit,best=L)  
tree.pred<-predict(mytree,newdata=test)  
sqrt(mean((tree.pred-test$medv)^2))
```

```
## [1] 5.134981
```


# Regression tree



# Regression tree - Boston data set



# Regression tree - Boston data set



# Regression tree - Boston data set


We can build a tree with all variable dataset (the syntax is similar to *lm()* function)

```
setup<-tree.control(nobs=nrow(train),  
                    mincut=2,minsize=4,mindev=0.005)  
tree.fit<-tree(medv~.,data=train,control=setup)  
plot(tree.fit)  
text(tree.fit,pretty = 0)
```

- A large tree is grown until a stopping criteria is reached.
- Function *cv.tree()* is used to perform cross-validation.
- The tree with the lowest cross validation error will be chosen.

```
cv.tree1<-cv.tree(tree.fit,K=5)  
plot(cv.tree1$size,cv.tree1$dev,type="b")
```

# Regression tree - Boston data set





# Regression tree - Boston data set

```
setup<-tree.control(nobs=nrow(train),  
                    mincut=2,minsize=4,mindev=0.005)  
tree.fit<-tree(medv~.,data=train,control=setup)  
L<-cv.tree1$size[which.min(cv.tree1$dev)]  
# number of leaves by cross validation  
L
```

```
## [1] 11
```

```
mytree<-prune.tree(tree.fit,best=L)  
tree.pred<-predict(mytree,newdata=test)  
sqrt(mean((tree.pred-test$medv)^2))
```

```
## [1] 4.5396
```

# Classification tree


- When the response variable  $Y$  is categorical we are in the classification setting.
- Growing a classification tree is quite similar to the task of growing a regression tree.
- In the classification setting, RSS cannot be used as a criterion for making the binary splits. An alternative to RSS is the classification error rate: At a node  $m$ , the *error rate* is calculated as follow

$$E = 1 - \max(p_{mk})$$

where  $p_{mk}$  is the proportion of training observations in the node that are from  $k^{th}$  class,  $k = 1, 2, \dots, K$ .

# Classification tree

Calculate the error rate in each node



# Classification tree

Calculate the error rate in each node



- Node  $m_1$ : we have  $p_1 = 1/2$ ;  $p_2 = 3/10$  and  $p_3 = 2/10$

$$E(m_1) = 1 - \max\{1/2, 3/10, 2/10\} = 0.5$$

- Node  $m_2$ : we have  $p_1 = 1/2$ ;  $p_2 = 1/2$  and  $p_3 = 0$

$$E(m_2) = 1 - \max\{1/2, 1/2, 0\} = 0.5$$

→ the error rate is not sensitive for growing a tree.

# Classification tree

- In practice, two other measures are preferable: *Gini index* and *Entropy*.

$$Gini\ index = \sum_{i=1}^K p_{mi} \times (1 - p_{mi})$$

$$Entropy = - \sum_{i=1}^K p_{mi} \log(p_{mi})$$

- The error rate, the Gini index, and the Entropy are measures of node purity i.e. these measures will take on a small value if the node is pure.
- The Gini index and the entropy are quite similar numerically.
- In a classification tree, either the Gini index or the entropy are used to evaluate the quality of a split.

# Classification tree

Calculate the Gini index and the entropy in each node



# Classification tree

Calculate the Gini index and the entropy in each node



- Node  $m_1$ : we have  $p_1 = 1/2$ ;  $p_2 = 3/10$  and  $p_3 = 2/10$

$$G(m_1) = 1/2 \times 1/2 + 3/10 \times 7/10 + 2/10 \times 8/10 = 0.62$$

$$D(m_1) = -(1/2 \log(1/2) + 3/10 \log(3/10) + 2/10 \log(2/10)) = 1.030$$

- Node  $m_2$ : we have  $p_1 = 1/2$ ;  $p_2 = 1/2$  and  $p_3 = 0$

$$G(m_2) = 1/2 \times 1/2 + 3/10 \times 7/10 + 2/10 \times 8/10 = 0.5$$

$$D(m_2) = -(1/2 \log(1/2) + 3/10 \log(3/10) + 2/10 \log(2/10)) = 0.693$$

# Classification tree

Write a function to calculate the error rate, the Gini index and the Entropy of a categorical variable.

```
purity<-function(v){  
  x<-as.numeric(v)  
  m<-min(x)  
  M<-max(x)  
  n<-length(x)  
  p<-rep(0,M-m+1)  
  for (j in m:M){  
    p[j]<-sum(x==j)/length(v)  
  }  
  result<-list()  
  result$err<-1-max(p)  
  result$gini<-sum(p*(1-p))  
  result$entropy<-sum(-p*log(p))  
  purity<-result
```



# Classification tree

Load **Default** data set from ISRL package, split data into training - test set (50% - 50%) and build a tree model where *default* variable depends on *balance* variable.

- Build a tree with two leaves?
- Which measure that function *tree* uses to split tree nodes?

```
dat<-Default
standardize<-function(x){x<-(x-mean(x,na.rm=TRUE))/sd(x,na.rm=TRUE)}
for (col in names(dat)){
  if(class(dat[,col]) %in% c("integer","numeric")){
    dat[,col]<-standardize(dat[,col])
  }
}

set.seed(1)
test_index<-createDataPartition(dat$default, times = 1, p = 0.5)
```


# Classification tree

```
K<-500
N<-dim(train)[1]
c1<-seq(min(train$balance)*0.99,max(train$balance)*0.99,length
Error<-rep(0,length(c1))
for(i in 1:K){
  leaf1<-train$default[train$balance<=c1[i]]
  leaf2<-train$default[train$balance>c1[i]]
  Error[i]<-length(leaf1)/N*purity(leaf1)$entropy+length(leaf2)
}

c1[which.min(Error)]
plot(c1,Error)
```

# Classification tree

## [1] 1.344704



# Classification tree

Build a tree model where *default* variable depends on other variables.  
Evaluate the performance of this model.

```
setup<-tree.control(nobs=nrow(train),  
                    mincut=2,minsize=4,mindev=0.005)  
tree.fit<-tree(default~.,data=train,control = setup)  
cv.tree1<-cv.tree(tree.fit,K=5)  
L<-cv.tree1$size[which.min(cv.tree1$dev)]  
# number of leaves by cross validation  
L
```

```
## [1] 6
```

```
mytree<-prune.tree(tree.fit,best=L)  
tree.pred<-predict(mytree,newdata=test,type="class")  
table(tree.pred,test$default)
```

```
##
```

# Decision tree versus linear model

In general, linear models assume the following form of  $f$ :

$$f(\mathbf{x}) = \beta_0 + \sum_{i=1}^p \beta_i x_i$$

whereas tree models assume a function of the form

$$f(\mathbf{x}) = \sum_{k=1}^M c_k \times \mathbb{I}_{\mathbf{x} \in R_k}$$

- Linear models are high-bias and low-variance method while tree models are low-bias and high-variance method
- If the relationship between the features and the response is approximated by a linear model, then linear models will outperform tree models.
- If there is a non-linear and complex relationship between the features and the response, then decision trees may outperform classical

# Advantages and disadvantages of tree-based models

- Trees are very easy to explain to people (even easier to explain than linear regression).
- Decision trees more closely mirror human decision-making than other approaches.
- Trees can be displayed graphically, and are easily interpreted even by a non-expert.
- Trees can easily handle qualitative predictors.
- Trees generally do not have the same level of predictive accuracy as other methods.
- Trees can be very non-robust i.e. a small change in the data can cause a large change in the final estimated tree (high variance)

Bagging, random forests, and boosting use trees to construct more powerful prediction models.